

Scientific computing using Cython: Best of both worlds!

March 18, 2017
Presenter: Simmi Mourya
Portland State University

Demo
github.com/simmimourya1/fossasia_17

Motivation




Illustration borrowed from: Stéfan van der Walt's presentation at Advanced Python Summer School, Kiel 2012.

Cython by example

Python

1x

```
def fib(n):  
    a,b = 1,1  
    for i in range(n):  
        a, b = a+b, a  
    return a
```

C/C++

100x

```
int fib(int n)  
{  
    int tmp, i, a, b;  
    a = b = 1;  
    for(i=0; i<n; i++){  
        tmp = a;  
        a += b;  
        b = tmp;  
    }  
}
```

Cython

80x

```
def fib(int n):  
    cdef int i, a, b  
    a, b = 1,1  
    for i in range(n):  
        a, b = a+b, a  
    return a
```

Cython in wild

Project	Cython files	Cython SLOC
sage	761	477,000
numpy	14	5,000
scipy	28	24,000
pandas	21	27,000
lxml	12	22,000
scikits-learn	35	15,000
scikits-image	48	11,000
mpi4py	48	12,000
yt	45	18,000

Projects master branches as of November 2014

Cython at Glance

Cython is used for compiling Python-like code to machine-code

- Supports a big subset of the Python language
- Conditions and loops run 2-8x faster, overall 30% faster for plain Python code (vs. Py2.5, using Py Bench)


Cython at Glance

In addition:

- Add types for speedups (hundreds of times)
- Easily use native libraries (C/C++/Fortran) directly

How it works:

- Cython code is turned into C code which uses the CPython API and runtime



Coding in Cython is
like coding in Python
and C at the same
time!

Use Case 1: Library Wrapping

- Cython is a popular choice for writing Python interface modules for C libraries

Use Case 2: Performance-critical code

- Python
- High-level
- Slow
- No variables typed
- C/C++/Fortran
- Lower-level
- Fast
- All variables typed

Common procedure: Where speed is needed, use a compiled language, then wrap the code for use from Python.

Breaking out of the global interpreter lock

- As a few of you might know, C Python has an infamous Global Interpreter Lock (GIL)

```
>>> import that
The Unwritten Rules of Python
1. You do not talk about the GIL.
2. You do NOT talk about the GIL.
3. Don't even mention the GIL. No seriously. ...
```

- It limits thread performance

Watch David Beazley' "Understanding the Python GIL" lecture.

Consider this simple script, that runs twice, sequentially, a `busy_sleep`, i.e. a function that simulate a CPU intensive task:

```
def busy_sleep( ):
    while n > 0:
        n -= 1
N = 99999999
busy_sleep(N)
    busy_sleep(N)
```

This takes 6.7 seconds to run.

Now consider the threaded version:

```
from threading import Thread
def busy_sleep(n):
    while n > 0:
        n -= 1
N = 99999999
t1 = Thread(target=busy_sleep, args=(N, ))
t2 = Thread(target=busy_sleep, args=(N, ))
t1.start()
t2.start()
t1.join()
t2.join()
```

This takes 11.1 seconds to run. What!?

Cython offers a wonderful context manager to run instructions without the GIL: *with nogil*.

Demo

Building Cython code

Building Cython code

Cython code must, unlike Python, be compiled. This happens in two stages:

- A .pyx file is compiled by Cython to a .c file, containing the code of a Python extension module
- The .c file is compiled by a C compiler into a .so file (or .pyd on Windows) which can be imported directly into a Python session.

Building Cython code

There are several ways to build Cython code:

- Write a distutils setup.py.
- Use pyximport, importing Cython .pyx files as if they were .py files (using distutils to compile and build in the background).
- Use the Jupyter notebook or the Sage notebook, both of which allow Cython code inline.

Building Cython modules using distutils

Imagine a simple “hello world” script in a file `hello.pyx`:

```
def say_hello_to(name):  
    print("Hello %s!" % name)
```

The following could be a corresponding `setup.py` script:

```
from distutils.core import setup  
from Cython.Build import cythonize  
  
setup(  
    name = 'Hello world app',  
    ext_modules = cythonize("hello.pyx"),  
)
```

Building Cython modules using distutils

To use this to build your Cython file use the command line options:

```
$ python setup.py build_ext --inplace
```

Which will leave a file in your local directory called helloworld.so (unix) or helloworld.pyd (Windows).


Now to use this file: start the python interpreter and simply import it as if it was a regular python module:

```
>>> import helloworld  
Hello World
```

pyximport: Cython Compilation the Easy Way


To load .pyx files directly on import, without having to write a setup.py file.

```
>>> import pyximport; pyximport.install()  
>>> import hello  
Hello World
```



Demo: Building
Cython modules
using Jupyter
notebook.

Language	Function call	Time (ms)	Speed, Python = 1
Python	<code>Fibo.fib(30)</code>	390	x 1
Cython	<code>cyFibo.fib(30)</code>	215	x 1.8
Cython	<code>cyFibo.fib_int(30)</code>	154	x 2.5
Cython	<code>cyFibo.fib_cdef(30)</code>	5.38	x72
Cython	<code>cyFibo.fib_cpdef(30)</code>	32.5	x12
C	<code>cFibo.fib(30)</code>	5.31	x73



The conclusions that I draw from this are:

1. Naive Cython does speed things up, but not by much (x1.8).
2. Optimised Cython is fairly effortless (in this case) and worthwhile (x2.5).
3. cdef is really valuable (x72).
4. cpdef gives a good improvement over def because the recursive case exploits C functions.
5. Cython's cdef is insignificantly different from the more complicated C extension that is our best attempt.

NumPy and Cython

- Cython provides fast access to NumPy arrays

MemoryViews

Demo

PROJECT DEMO
CYVLFEAT
GOOGLE SUMMER OF CODE 2016
@ PORTLAND STATE UNIVERSITY

A CYTHON / PYTHON WRAPPER FOR VLFEAT



WHAT IS VLFEAT?



WHAT IS CYVLFEAT?

[GITHUB.COM/MENPO/CYVLFEAT](https://github.com/menpo/cyvlfeat)

A photograph of a man and two children playing soccer on a green field. The man, in the center, is wearing a white soccer jersey with red accents and red shorts. He is smiling and has his arms around two young boys. The boy on the right is wearing a black soccer jersey and black shorts, also smiling. The boy on the left is wearing a black hoodie and black shorts, and is walking away. The background shows a black fence and some buildings under a cloudy sky.

CONTRIBUTIONS ARE WELCOME!
[GITHUB.COM/SIMMIMOURYAL/CYVLFEAT](https://github.com/simmimouryal/cyvlfeat)

A dark, grainy background image showing a soccer game at night. A bright stadium light illuminates the center of the field, creating a circular glow. Several players in white and dark jerseys are visible on the field, some in motion. The overall scene is dimly lit, with the primary light source being the central stadium light.

FEATURES WHICH NEED MORE WORK

1. BINSUM
2. KDTREE MODULE

QUESTIONS ?



THANK YOU

